

```

!pip install -U scipy==1.2.0
from scipy.optimize import fmin_l_bfgs_b
from scipy.misc import imsave
import warnings
warnings.filterwarnings('ignore')

%matplotlib inline
import importlib
import tensorflow as tf
import tensorflow.keras as K
import math, datetime, pandas as pd, numpy as np
import matplotlib.pyplot as plt, random, pickle, glob, os
import sklearn
from PIL import Image
import tarfile
import cv2
import random

```

```

Collecting scipy==1.2.0
  Downloading scipy-1.2.0-cp37-cp37m-manylinux1_x86_64.whl (26.6 MB)
    |████████████████████████████████████████| 26.6 MB 75 kB/s
Requirement already satisfied: numpy>=1.8.2 in /usr/local/lib/python3.7/dist-packages
Installing collected packages: scipy
  Attempting uninstall: scipy
    Found existing installation: scipy 1.4.1
    Uninstalling scipy-1.4.1:
      Successfully uninstalled scipy-1.4.1
ERROR: pip's dependency resolver does not currently take into account all the packages
jax 0.2.21 requires scipy>=1.2.1, but you have scipy 1.2.0 which is incompatible.
albumations 0.1.12 requires imgaug<0.2.7,>=0.2.5, but you have imgaug 0.2.9 which
Successfully installed scipy-1.2.0

```

## ▼ Loading data

```

input_size=240

path='/content/'

img_list = glob.glob(path+'**/*.jpg', recursive=True)
n = len(img_list); n

1

train_df=pd.read_csv('train_data.csv')
train_df.shape

(1, 3)

train_list=train_df.ids

```

```
len(train_list)
```

```
1
```

```
val_list = glob.glob('/content/*.jpg*', recursive=True)
len(val_list)
```

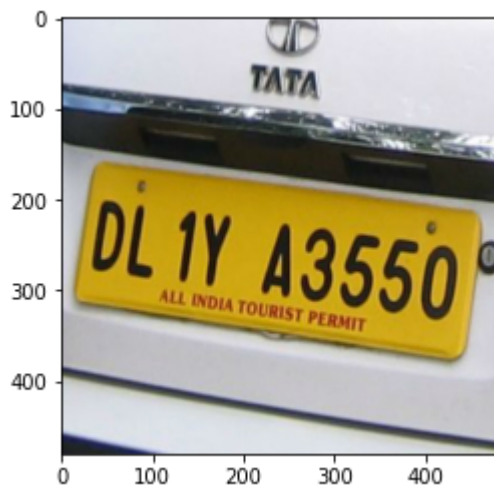
```
1
```

```
def read_image(img_path,scale=2):
    global input_size
    img = cv2.imread(img_path, cv2.IMREAD_COLOR)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img_lr = cv2.resize(img, (input_size,input_size),interpolation=cv2.INTER_CUBIC)
    img_hr = cv2.resize(img, (input_size*scale,input_size*scale),interpolation=cv2.INTER_C
    return img_lr,img_hr
```

```
_,img=read_image(train_list[0])
print(img.shape)
plt.imshow(img)
```

```
(480, 480, 3)
```

```
<matplotlib.image.AxesImage at 0x7f53aab51a50>
```



```
_,img=read_image(val_list[0])
print(img.shape)
plt.imshow(img)
```

```
(480, 480, 3)
<matplotlib.image.AxesImage at 0x7f53a963f950>
```



## ▼ Super Resolution Model



```
def conv_block(x, filters, kernel, stride=(1,1), mode='same', act=True):
    x = K.layers.Conv2D(filters, kernel_size=kernel, strides=stride, padding=mode)(x)
    return K.layers.Activation('relu')(x) if act else x
```



```
def res_block(ip,nf=16):
    x = conv_block(ip, nf, 3, (1,1))
    x = conv_block(x, nf, 3, (1,1), act=False)
    return K.layers.Add()([x,ip])
```

```
def up_block(x,nf):
    x = K.layers.UpSampling2D()(x)
    x = conv_block(x,nf,kernel=(1,1))
    return x
```

```
def get_srmodel(shape=(None,None,3)):
    inp=K.Input(shape)

    x=conv_block(inp, 16, 3, (1,1))
    for i in range(5): x=res_block(x,16)
```

```
    x=up_block(x,nf=32)
    #    for i in range(2): x=res_block(x,32)

    #    x=up_block(x,nf=64)
    #    for i in range(1): x=res_block(x,64)
```

```
    x = K.layers.Conv2D(3,(3,3),padding='same')(x)
    return inp,x
```

```
inp,output=get_srmodel((None,None,3))
sr_model=K.Model(inp,output)
sr_model.summary(110)
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected
input_1 (InputLayer)	[(None, None, None, 3)]	0	
conv2d (Conv2D)	(None, None, None, 16)	448	input_1[0]
activation (Activation)	(None, None, None, 16)	0	conv2d[0]

conv2d_1 (Conv2D)	(None, None, None, 16)	2320	activation_1
activation_1 (Activation)	(None, None, None, 16)	0	conv2d_1[0][0]
conv2d_2 (Conv2D)	(None, None, None, 16)	2320	activation_2
add (Add)	(None, None, None, 16)	0	conv2d_2[0][0]
conv2d_3 (Conv2D)	(None, None, None, 16)	2320	add[0][0]
activation_2 (Activation)	(None, None, None, 16)	0	conv2d_3[0][0]
conv2d_4 (Conv2D)	(None, None, None, 16)	2320	activation_3
add_1 (Add)	(None, None, None, 16)	0	conv2d_4[0][0]
conv2d_5 (Conv2D)	(None, None, None, 16)	2320	add_1[0][0]
activation_3 (Activation)	(None, None, None, 16)	0	conv2d_5[0][0]
conv2d_6 (Conv2D)	(None, None, None, 16)	2320	activation_4
add_2 (Add)	(None, None, None, 16)	0	conv2d_6[0][0]
conv2d_7 (Conv2D)	(None, None, None, 16)	2320	add_2[0][0]
activation_4 (Activation)	(None, None, None, 16)	0	conv2d_7[0][0]
conv2d_8 (Conv2D)	(None, None, None, 16)	2320	activation_5
add_3 (Add)	(None, None, None, 16)	0	conv2d_8[0][0]
conv2d_9 (Conv2D)	(None, None, None, 16)	2320	add_3[0][0]
activation_5 (Activation)	(None, None, None, 16)	0	conv2d_9[0][0]
conv2d_10 (Conv2D)	(None, None, None, 16)	2320	activation_6
add_4 (Add)	(None, None, None, 16)	0	conv2d_10[0][0]
up_sampling2d (UpSampling2D)	(None, None, None, 16)	0	add_4[0][0]
conv2d_11 (Conv2D)	(None, None, None, 32)	544	up_sampling2d[0][0]

## ▼ loading vgg for calculating perceptual loss from one of its layer

```

vgg_inp=K.Input(shape=outp.shape[1:])
vgg= K.applications.VGG16(include_top=False,
                           input_tensor=vgg_inp)

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_data\\_format.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_data_format.h5)  
58892288/58889256 [=====] - 0s 0us/step

58900480/58889256 [=====] - 0s 0us/step

```
for l in vgg.layers: l.trainable=False
```

```
# preproc_layer = K.layers.Lambda(preproc)
```

```
#Here we are using vgg layer at index 36 to be the layer to calculate loss between Target
vgg_out_layer = vgg.get_layer(index=5).output
```

```
# making model Model(inputs, outputs)
vgg_content = K.Model(vgg_inp, vgg_out_layer)
```

```
vgg_content.summary(110)
```

```
Model: "model_1"
```

Layer (type)	Output Shape
input_2 (InputLayer)	[(None, None, None, 3)]
block1_conv1 (Conv2D)	(None, None, None, 64)
block1_conv2 (Conv2D)	(None, None, None, 64)
block1_pool (MaxPooling2D)	(None, None, None, 64)
block2_conv1 (Conv2D)	(None, None, None, 128)
block2_conv2 (Conv2D)	(None, None, None, 128)
Total params: 260,160	
Trainable params: 0	
Non-trainable params: 260,160	

## ▼ Data Generator and Metrics

```
def randomHorizontalFlip(img, u=0.5):
    if np.random.random() < u:
        img = cv2.flip(img, 1)
    return img
def randomVerticalFlip(img, u=0.5):
    if np.random.random() < u:
        img = cv2.flip(img, 0)
    return img
```

```
# def randomCrop(img):
#     global input_size
#     h=input_size*2
```

```

#     assert img.shape[0] >= h
#     assert img.shape[1] >= h
#     x = random.randint(0, img.shape[1] - h)
#     y = random.randint(0, img.shape[0] - h)
#     img = img[y:y+h, x:x+h]
#     return img

def train_generator():
    global batch_size
    global input_size
    while True:
        for start in range(0, len(train_list), batch_size):
            x_batch = []
            y_batch = []
            end = min(start + batch_size, len(train_list))
            ids_train_batch = train_list[start:end]
            for i,ids in enumerate(ids_train_batch):
                img = cv2.imread(ids)
                img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                tar = cv2.resize(img, (input_size*2, input_size*2),interpolation=cv2.INTER_LINEAR)
                tar = randomHorizontalFlip(tar)
                tar = randomVerticalFlip(tar)
                img1 = cv2.resize(tar, (input_size, input_size),interpolation=cv2.INTER_CUBIC)
                x_batch.append(img1)
                y_batch.append(tar)
            x_batch = np.array(x_batch, np.float32) / 255.
            y_batch = np.array(y_batch, np.float32) / 255.
            yield x_batch, y_batch

# i am using Set5 dataset for validation please download the data
def valid_generator():

    global batch_size
    batch_size=32
    global input_size
    while True:
        for start in range(0, len(val_list), batch_size):
            x_batch = []
            y_batch = []
            end = min(start + batch_size, len(val_list))
            ids_valid_batch = val_list[start:end]
            for i,ids in enumerate(ids_valid_batch):
                img = cv2.imread(ids)
                img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                img1 = cv2.resize(img, (input_size, input_size),interpolation=cv2.INTER_CUBIC)
                img2 = cv2.resize(img, (input_size*2, input_size*2),interpolation=cv2.INTER_LINEAR)
                x_batch.append(img1)
                y_batch.append(img2)
            x_batch = np.array(x_batch, np.float32) / 255.
            y_batch = np.array(y_batch, np.float32) / 255.
            yield x_batch, y_batch

l=next(valid_generator())

```

```
l=next(valid_generator())
```

```

def psnr(y_true,y_pred):
    return tf.image.psnr(y_true,y_pred,1.0)
def ssim(y_true,y_pred):
    return tf.image.ssim(y_true,y_pred,1.0)

# This is our perceptual loss function
def perceptual_loss(y_true,y_pred):
    # mse=K.losses.mean_squared_error(y_true,y_pred)
    y_t=vgg_content(y_true)
    y_p=vgg_content(y_pred)
    loss=K.losses.mean_squared_error(y_t,y_p)
    return loss

learning_rate=0.001
adam=K.optimizers.Adam(lr=learning_rate)
sr_model.compile(optimizer=adam,loss=perceptual_loss,metrics=[psnr,ssim])

batch_size=16
input_size=32

def fit(model,epoch=2):
    model.fit_generator(generator=train_generator(),
                        steps_per_epoch=np.ceil(float(len(train_list)) / float(batch_size)),
                        epochs=epoch,
                        verbose=1,
                        validation_data=valid_generator(),
                        shuffle=True,
                        validation_steps=np.ceil(float(len(val_list)) / float(batch_size)))
    return model

sr_model=fit(sr_model,10)

Epoch 1/10
1/1 [=====] - 3s 3s/step - loss: 40.1719 - psnr: 4.8845 - ss
Epoch 2/10
1/1 [=====] - 0s 167ms/step - loss: 36.9319 - psnr: 5.6557 -
Epoch 3/10
1/1 [=====] - 0s 177ms/step - loss: 33.4245 - psnr: 6.0533 -
Epoch 4/10
1/1 [=====] - 0s 194ms/step - loss: 31.6264 - psnr: 6.0062 -
Epoch 5/10
1/1 [=====] - 0s 170ms/step - loss: 26.9437 - psnr: 6.0627 -
Epoch 6/10
1/1 [=====] - 0s 174ms/step - loss: 24.0374 - psnr: 6.2598 -
Epoch 7/10
1/1 [=====] - 0s 176ms/step - loss: 24.9620 - psnr: 6.4472 -
Epoch 8/10
1/1 [=====] - 0s 187ms/step - loss: 22.2344 - psnr: 7.0352 -
Epoch 9/10
1/1 [=====] - 0s 171ms/step - loss: 22.4084 - psnr: 7.2959 -
Epoch 10/10
1/1 [=====] - 0s 181ms/step - loss: 21.2397 - psnr: 7.5057 -

```

## progressive resizing

```
# input_size=64
# sr_model=fit(sr_model,15)
```

```
# input_size=96
# sr_model=fit(sr_model,10)
```

```
# input_size=128
# sr_model=fit(sr_model,15)
```

```
input_size=160
sr_model=fit(sr_model,10)
```

```
Epoch 1/10
1/1 [=====] - 3s 3s/step - loss: 7.8641 - psnr: 8.2727 - ss
Epoch 2/10
1/1 [=====] - 3s 3s/step - loss: 7.4017 - psnr: 8.4020 - ss
Epoch 3/10
1/1 [=====] - 3s 3s/step - loss: 7.2271 - psnr: 8.6255 - ss
Epoch 4/10
1/1 [=====] - 3s 3s/step - loss: 6.6564 - psnr: 8.9619 - ss
Epoch 5/10
1/1 [=====] - 3s 3s/step - loss: 6.3465 - psnr: 9.6815 - ss
Epoch 6/10
1/1 [=====] - 3s 3s/step - loss: 5.9944 - psnr: 10.4393 - ss
Epoch 7/10
1/1 [=====] - 3s 3s/step - loss: 5.4895 - psnr: 11.3742 - ss
Epoch 8/10
1/1 [=====] - 3s 3s/step - loss: 5.2313 - psnr: 12.2730 - ss
Epoch 9/10
1/1 [=====] - 3s 3s/step - loss: 5.0863 - psnr: 12.9626 - ss
Epoch 10/10
1/1 [=====] - 3s 3s/step - loss: 4.8507 - psnr: 13.3464 - ss
```

```
input_size=224
sr_model=fit(sr_model,10)
```

```
Epoch 1/10
1/1 [=====] - 7s 7s/step - loss: 3.1857 - psnr: 13.8676 - ss
Epoch 2/10
1/1 [=====] - 6s 6s/step - loss: 3.0207 - psnr: 13.8468 - ss
Epoch 3/10
1/1 [=====] - 6s 6s/step - loss: 3.0525 - psnr: 13.5905 - ss
Epoch 4/10
1/1 [=====] - 6s 6s/step - loss: 2.9152 - psnr: 13.6006 - ss
Epoch 5/10
1/1 [=====] - 6s 6s/step - loss: 2.6122 - psnr: 13.6593 - ss
Epoch 6/10
1/1 [=====] - 6s 6s/step - loss: 2.6943 - psnr: 14.0638 - ss
Epoch 7/10
1/1 [=====] - 6s 6s/step - loss: 2.5818 - psnr: 14.4705 - ss
```



```
Epoch 8/10
1/1 [=====] - 6s 6s/step - loss: 2.3301 - psnr: 14.8441 - ss
Epoch 9/10
1/1 [=====] - 7s 7s/step - loss: 2.3803 - psnr: 15.3961 - ss
Epoch 10/10
1/1 [=====] - 6s 6s/step - loss: 2.4722 - psnr: 15.7467 - ss
```



```
input_size=160
```

```
tlr, hlr = read_image(val_list[0])
tlr = np.expand_dims(tlr, axis=0)
pred = sr_model.predict(tlr/255.)
```

```
import cv2
import numpy
from google.colab.patches import cv2_imshow
from matplotlib import pyplot as plt
plt.axis('off')
plt.imshow(pred[0])
plt.savefig('main.jpg')
fig, axs = plt.subplots(1, 3, figsize=(20, 20))
ax = axs.flat
```

```
ax[0].imshow(tlr[0]/255.)
ax[1].imshow(pred[0])
ax[2].imshow(hlr/255.)
```



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or  
 Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or  
 <matplotlib.image.AxesImage at 0x7f53a2931c90>



```
K.models.save_model(sr_model, 'sr_2x.h5')
```



```
import cv2
import numpy as np
from matplotlib import pyplot as plt
import cv2
from google.colab.patches import cv2_imshow
```

```
img = cv2.imread('/content/dem.jpg')
rows, cols, ch = img.shape
```

```
pts1 = np.float32([[50, 50],
                   [200, 50],
                   [50, 200]])
```

```
pts2 = np.float32([[10, 110],
                   [200, 100],
                   [0, 200]])
```

```
M = cv2.getAffineTransform(pts1, pts2)
dst = cv2.warpAffine(img, M, (cols, rows))
cv2.imwrite('tilted.jpg', dst)
plt.subplot(121)
plt.imshow(img)
plt.title('Input')
```

```
plt.subplot(122)
plt.imshow(dst)
plt.title('Output')
plt.show()
```

